


Data Structure:

B-Tree

m-way search tree

- Binary trees are not quite appropriate for data stored on disks
 - we assumed all data is kept in main memory
 - what if the data is kept in external disk?
 - disk access is much slower than memory access
 - disk is partitioned into blocks (pages) and the access time of a word is the same as that of the entire block containing the word
 - we need to reduce the number of disk access
 - make each node of the tree wider (m-way search tree)




m-way search tree

- In a tree of degree m and height h
 - the maximum number of **nodes** is $(m^h - 1) / (m - 1)$
 $(m^0 + m^1 + m^2 + \dots m^{h-1})$
 - the maximum number of **elements** in an m -way tree of **height h** is $m^h - 1$
(since each node has at most $m-1$ elements)
- a binary tree with $h=3$ has 7 elements in the tree
- a 200-way tree with $h=3$ has $200^3 - 1 = 8 * 10^6 - 1$ nodes

B-Tree

- a B-tree of **order m** is an m -way search tree with the following properties
 - the root has **at least 2 children**
 - each node has upto **$m-1$ keys**
 - all **external nodes** are at **the same level (perfectly balanced)**
 - all **internal nodes** (except the root) have between **$\lceil m/2 \rceil$ and m children**
 - when $m=3$, all internal nodes of B-tree have a degree of either 2 or 3 (2-3 tree)
 - when $m=4$, all internal nodes of B-tree have a degree of 2, 3, or 4 (2-3-4 tree)



B-Tree

- a B-tree of height h
- best case: the tree is splitting widely

$$n = m^h - 1$$

$$h = \lceil \log_m (n + 1) \rceil$$

$$\log_m n = \frac{\log n}{\log m} = O(\log n)$$

- worst case: the tree is splitting $\lceil m/2 \rceil$ ways

$$\log_{\lceil \frac{m}{2} \rceil} n = \frac{\log n}{\log \lceil \frac{m}{2} \rceil} = O(\log n)$$


B-Tree: node structure

```
#define order 3
```

```
struct B_node {  
    int  order;          /* number of children */  
    B_node *child[order]; /* children pointers */  
    int  key[order-1];    /* keys          */  
}
```

search

- When we arrive an internal node with key $k_1 < k_2, \dots < k_{m-1}$, **search for x** in this list (either linearly or by binary search)
 - if you found x , you are done
 - otherwise, find the index i such that $k_i < x < k_{i+1}$, and recursively search the subtree pointed by p_i .
- Complexity = $\log m \cdot \log_m n = O(\log n)$




insertion

- find the appropriate leaf into which the node can be inserted
 - if the leaf is not full ($< m-1$ keys), insert it

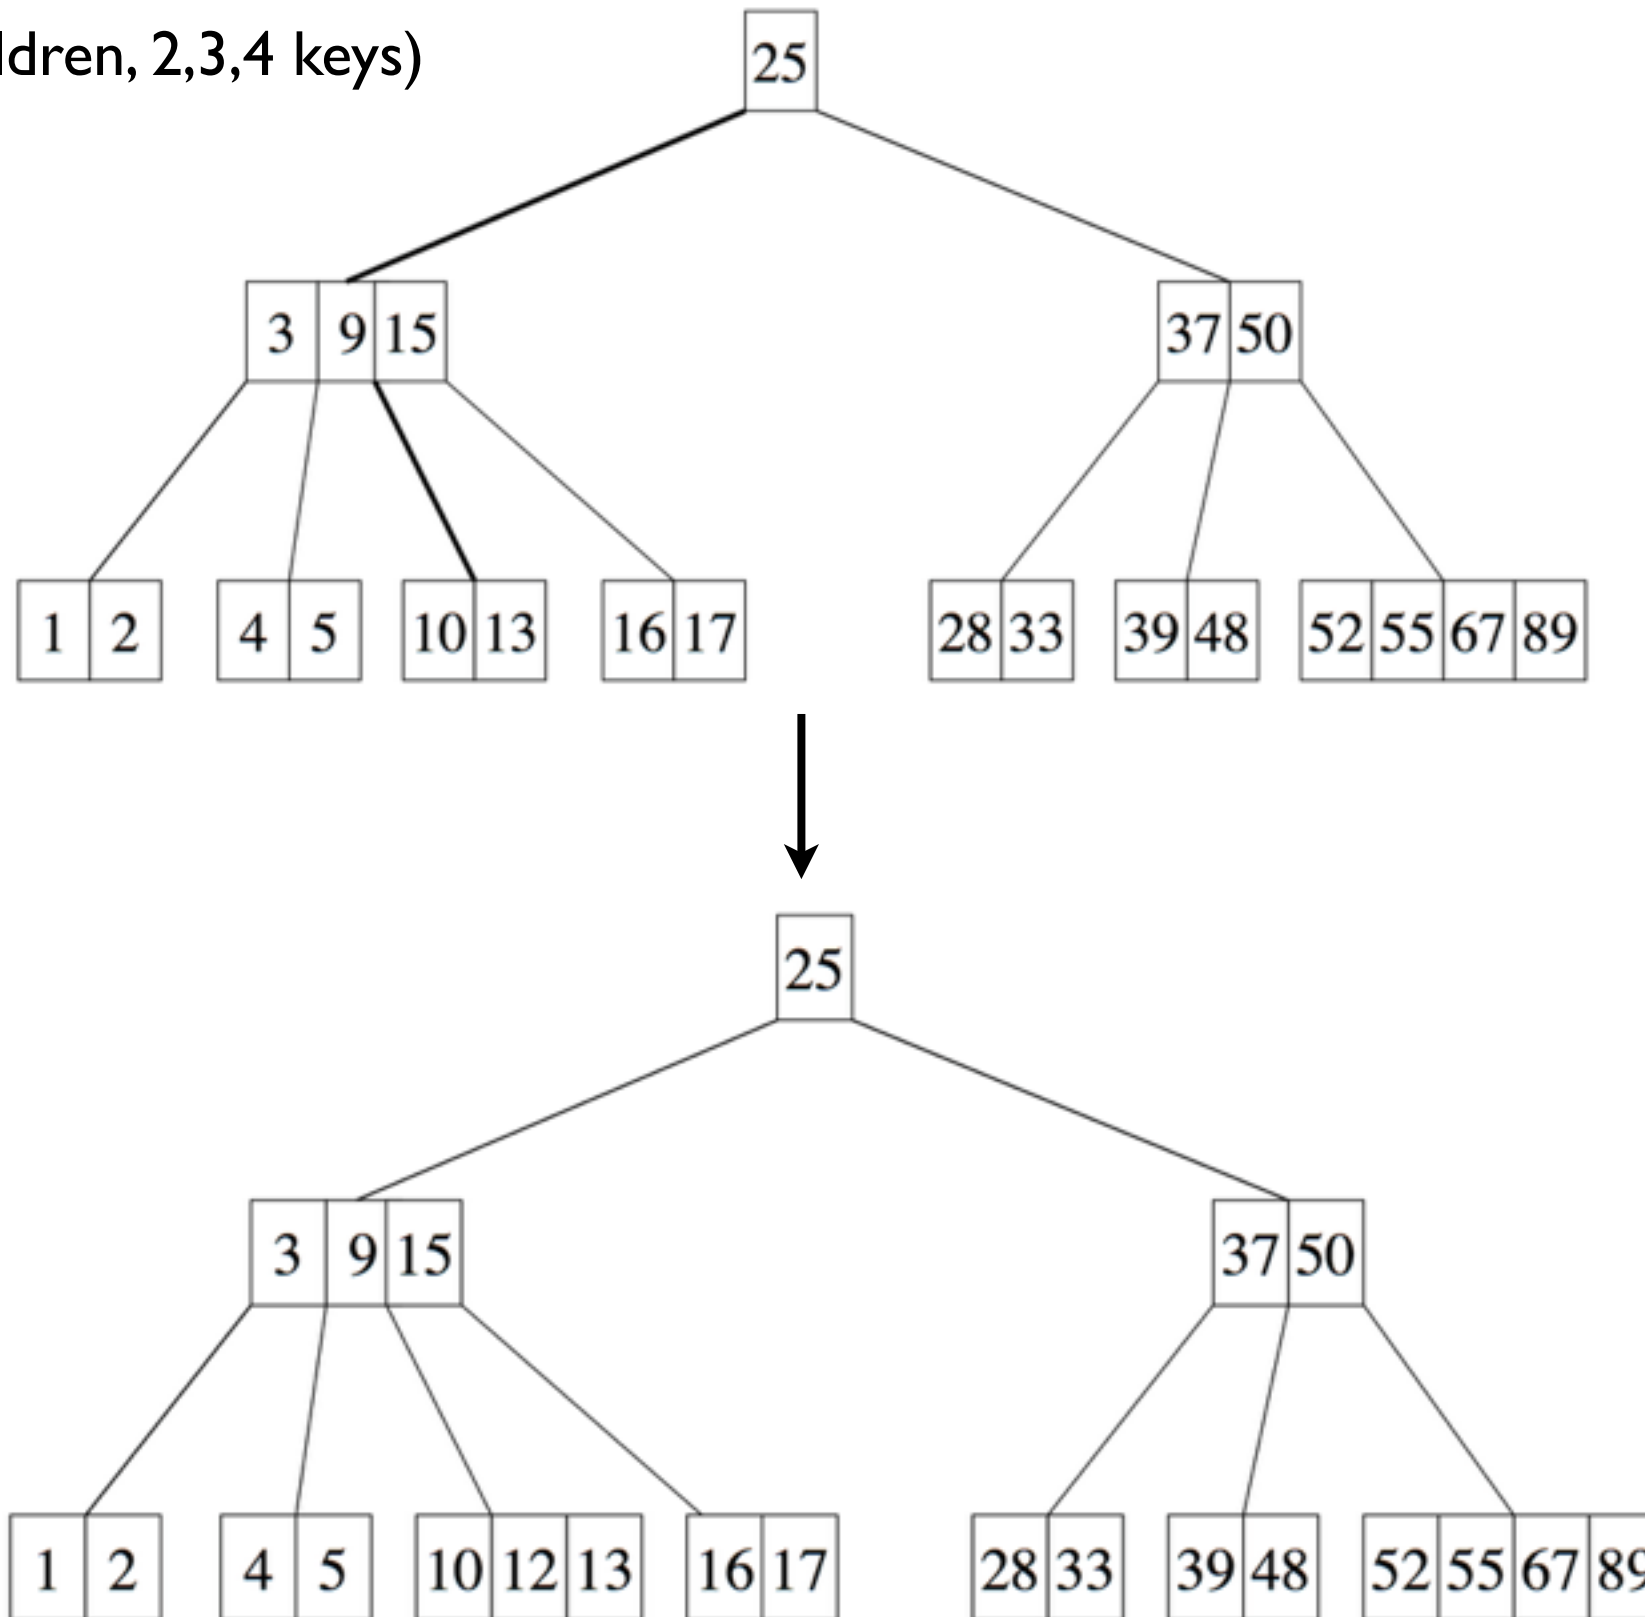
insertion

m= 5 (3,4,or 5 children, 2,3, or 4 keys)
insert 12



insertion

m= 5 (3,4,5 children, 2,3,4 keys)
insert 12




insertion

- find the appropriate leaf into which the node can be inserted
 - if the leaf is not full ($< m-1$ keys), insert it
 - if the node overflows, restore the balance
 - key rotation (if there is a space in the sibling node)
 - node split


insertion

- **key rotation**: check for siblings for rotation into the B-tree of $m=3$



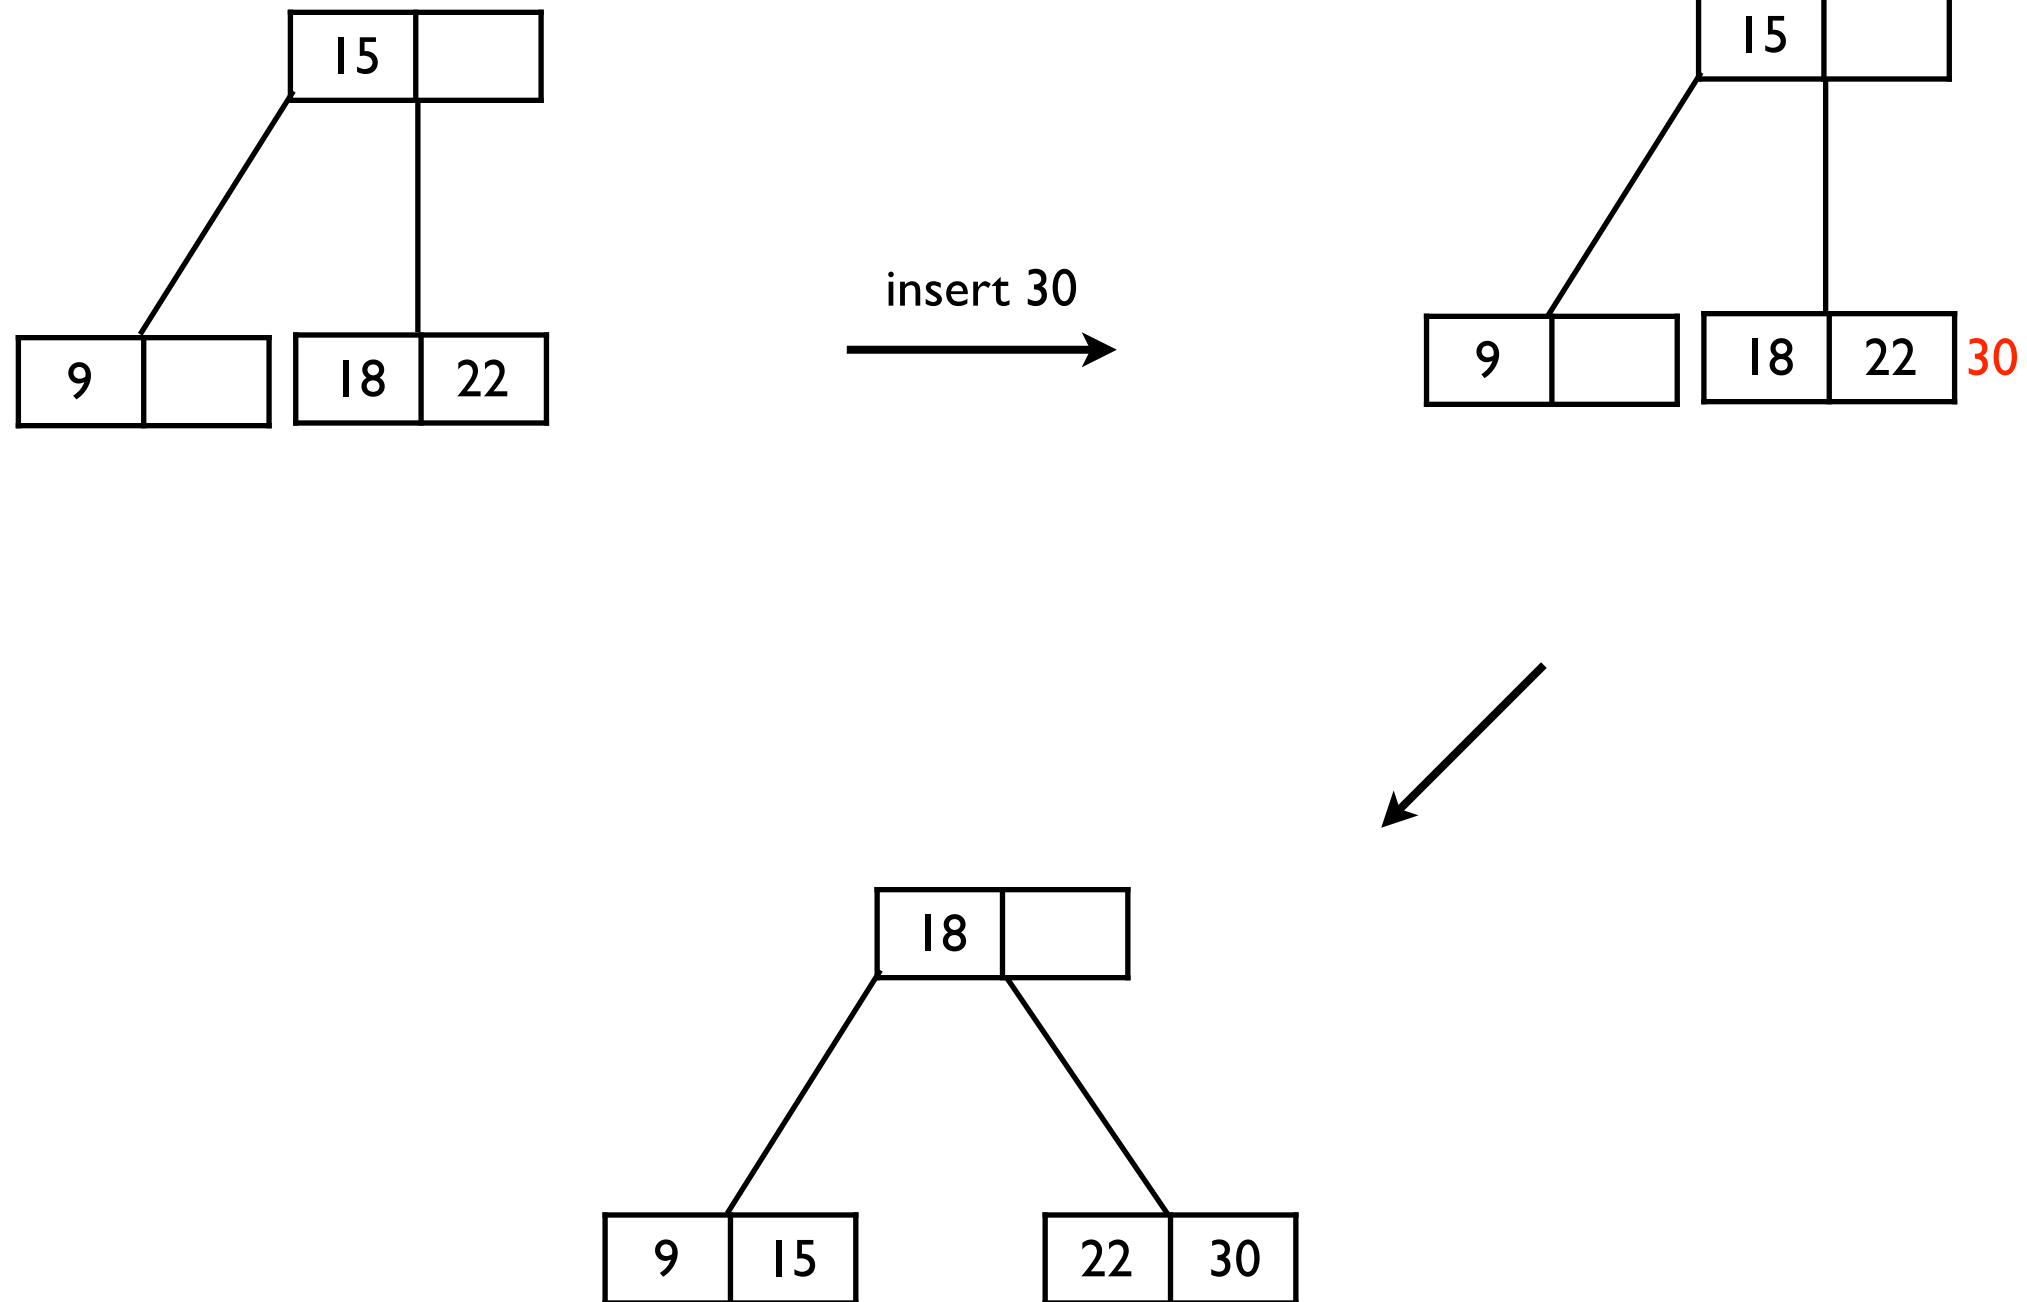
insertion

- **key rotation**: check for siblings for rotation into the B-tree of $m=3$




insertion

- **key rotation**: check for siblings for rotation into the B-tree of $m=3$




insertion




insert 60




insertion





insert 60



insertion



insert 60



insertion


- node split

- if we have a node with m keys after insertion (overflow), split the node into three groups
 - (a) a node with the keys smaller than the middle key
 - (b) a node with the middle key
 - (c) a node with the keys greater than the middle key
- make (a) and (c) as new nodes and push (b) to the parent
- if the parent overflows, repeat the process
- if the root overflows, create a new node with 2 children


insertion

m= 3 (2,3 children, 1,2 keys)

insert 30




insert 30




insertion

m= 3 (2,3 children, 1,2 keys)


insert 30



insert 30




find the middle one and
push it to the parent node




insertion

m= 3 (2,3 children, 1,2 keys)



insert 30



insert 30




find the middle one and
push it to the parent node




insertion

m= 5 (3, 4, 5 children, 2,3,4 keys)

insert 53






insert 53



insertion

m= 5 (3, 4, 5 children, 2,3,4 keys)





insert 53



insertion


m= 5 (3, 4, 5 children, 2,3,4 keys)

insert 53




insertion

m= 5 (3, 4, 5 children, 2,3,4 keys)



insertion


m= 5 (3, 4, 5 children, 2,3,4 keys)




insert 67

insertion

m= 5 (3, 4, 5 children, 2,3,4 keys)

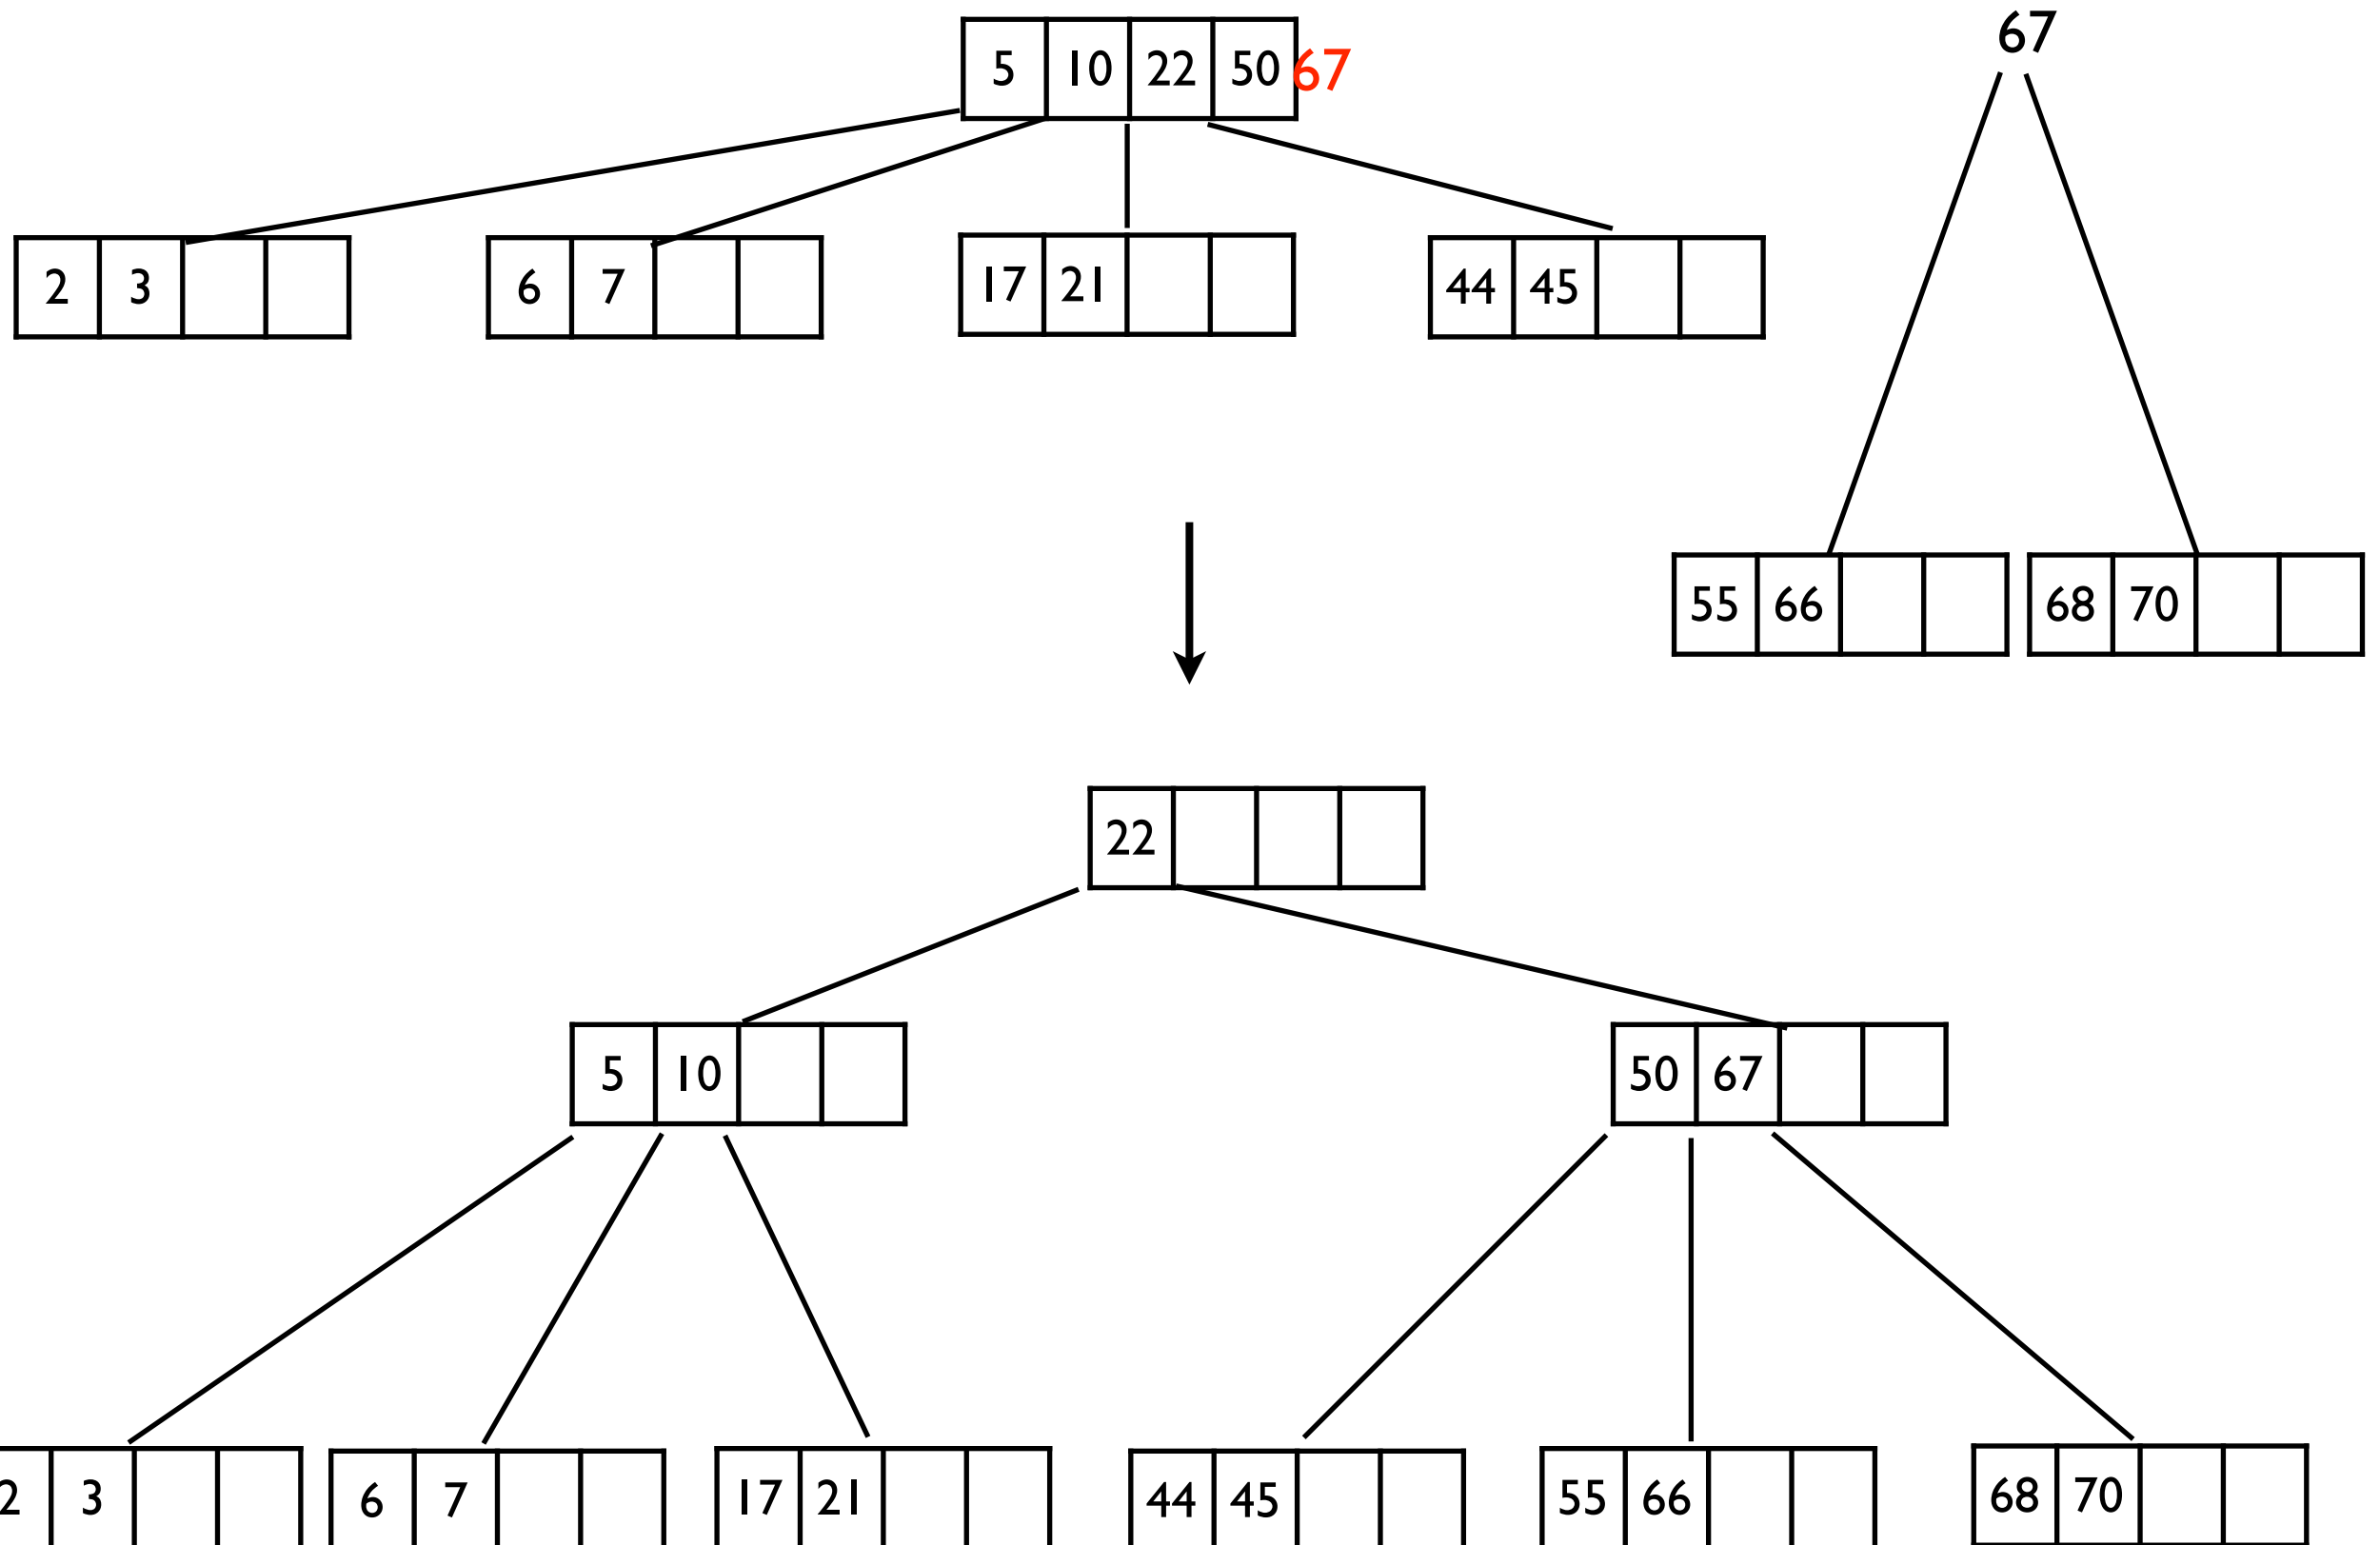


insert 67



insertion

m= 5 (3, 4, 5 children, 2,3,4 keys)







deletion

- find a suitable replacement which is the largest key in the left child (or the smallest in the right) and move it to fill the hole
 - key rotation
 - node merging


deletion

- **key rotation:** check for siblings for rotation

m= 3 (2,3 children, 1,2 keys)




deletion




delete 60




deletion




delete 60




deletion



delete 60





key rotation




deletion

■ node merging

- no sibling that can be rotated
- move down the intermediate node from the parent and put it in the new node
- if this might cause underflow in the parent node, repeat the process




deletion




delete 95



deletion




delete 95




deletion


m= 3 (2,3 children, 1,2 keys)



delete 95




node merging



deletion

m= 5 (3, 4, 5 children, 2,3,4 keys)

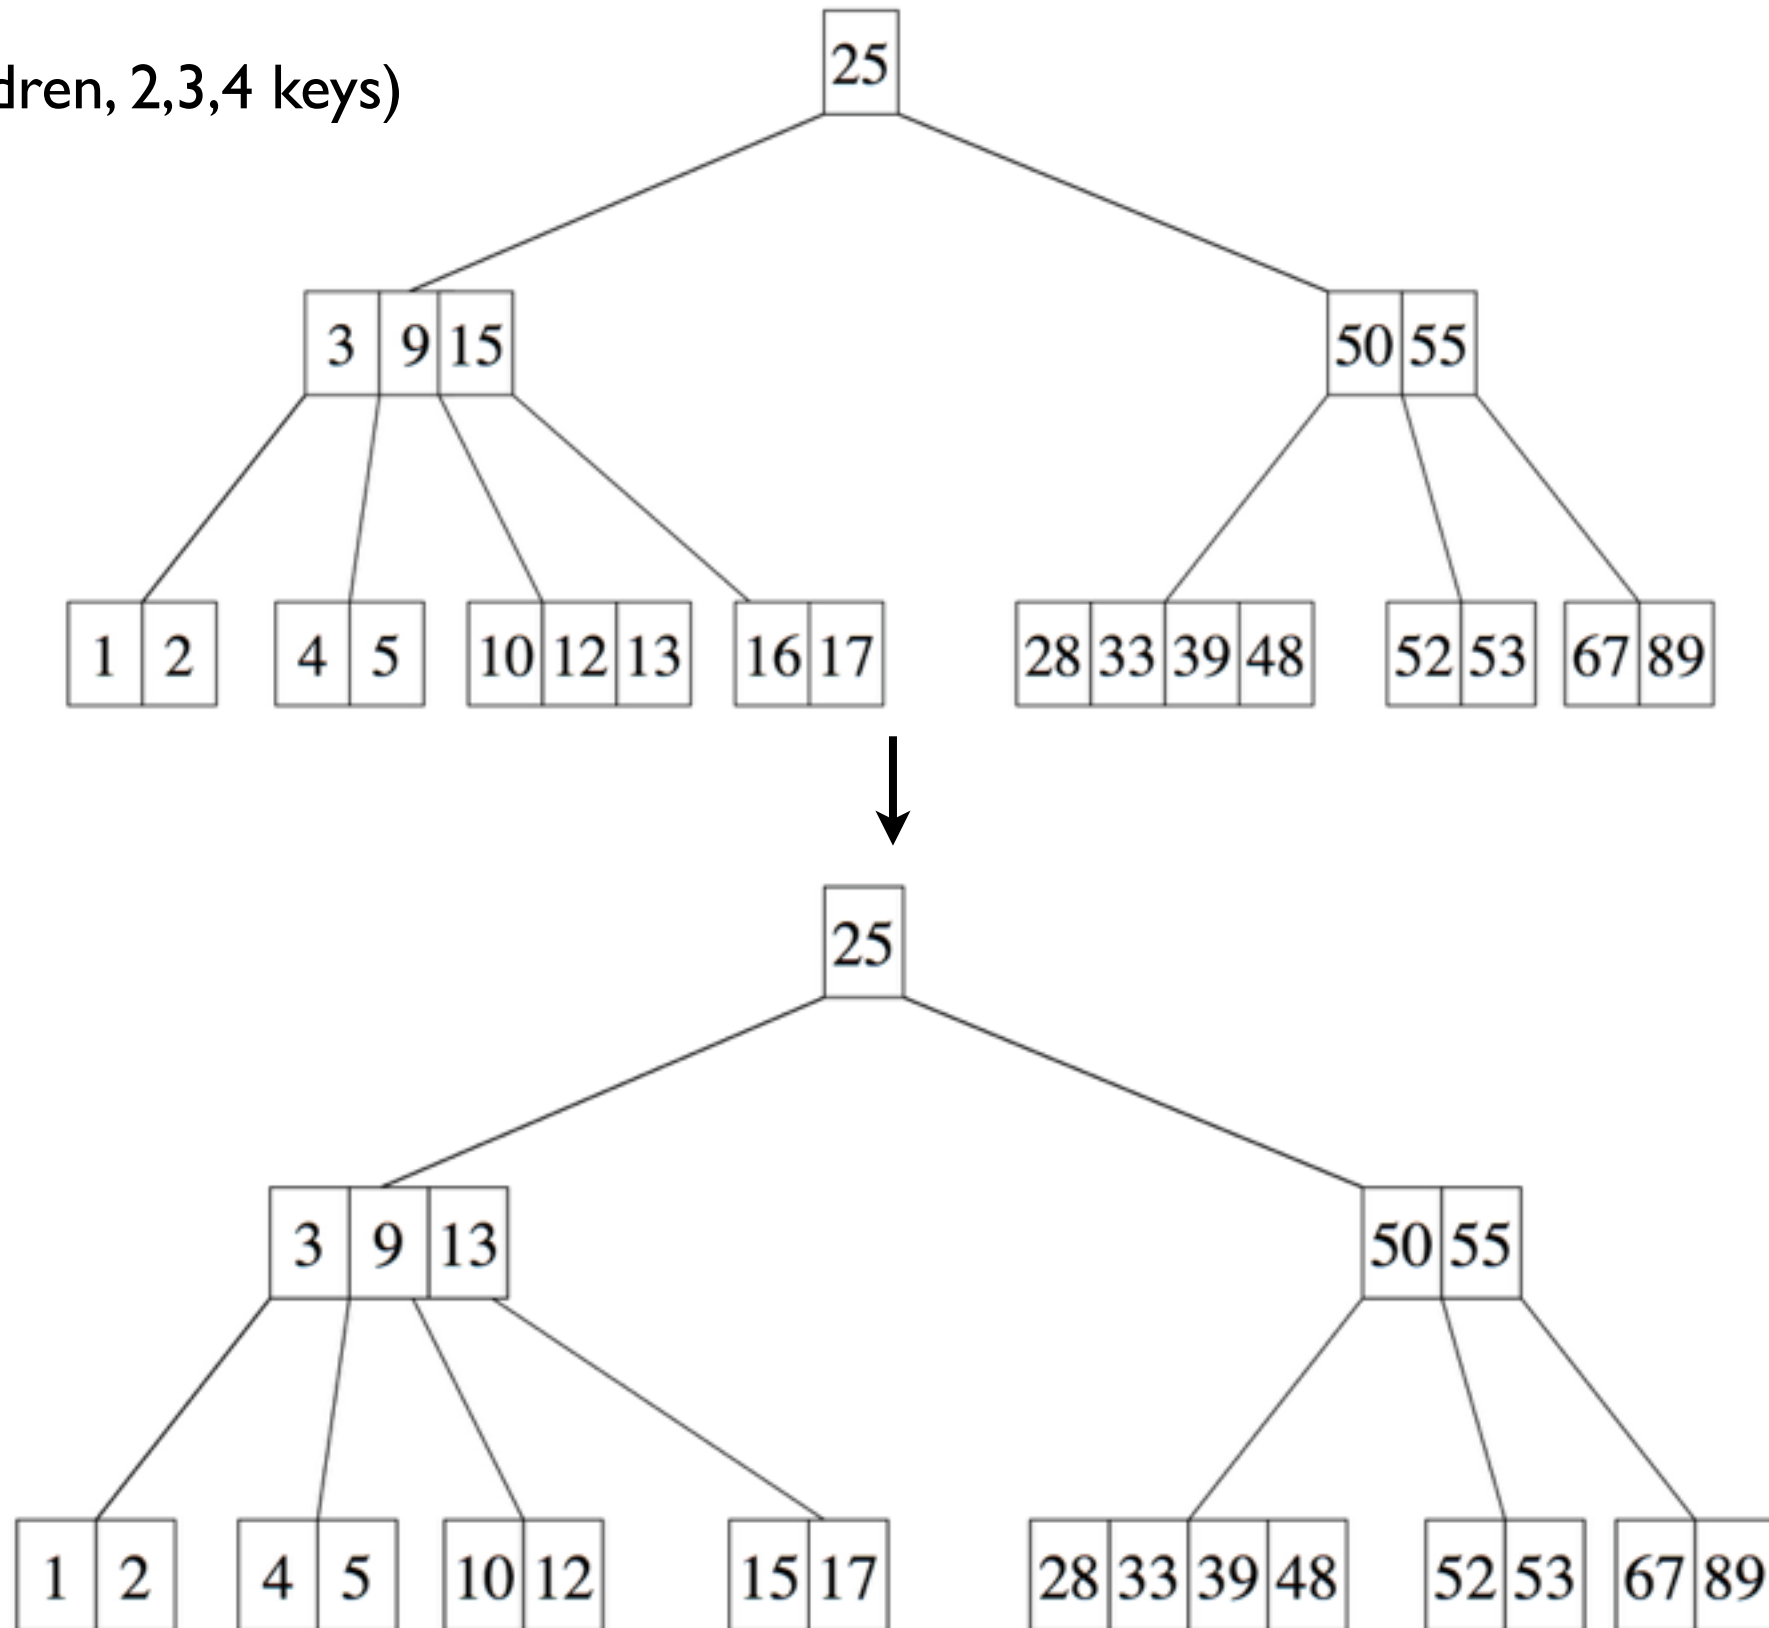
delete 16



deletion

m= 5 (3, 4, 5 children, 2,3,4 keys)


delete 16



deletion

m= 5 (3, 4, 5 children, 2,3,4 keys)

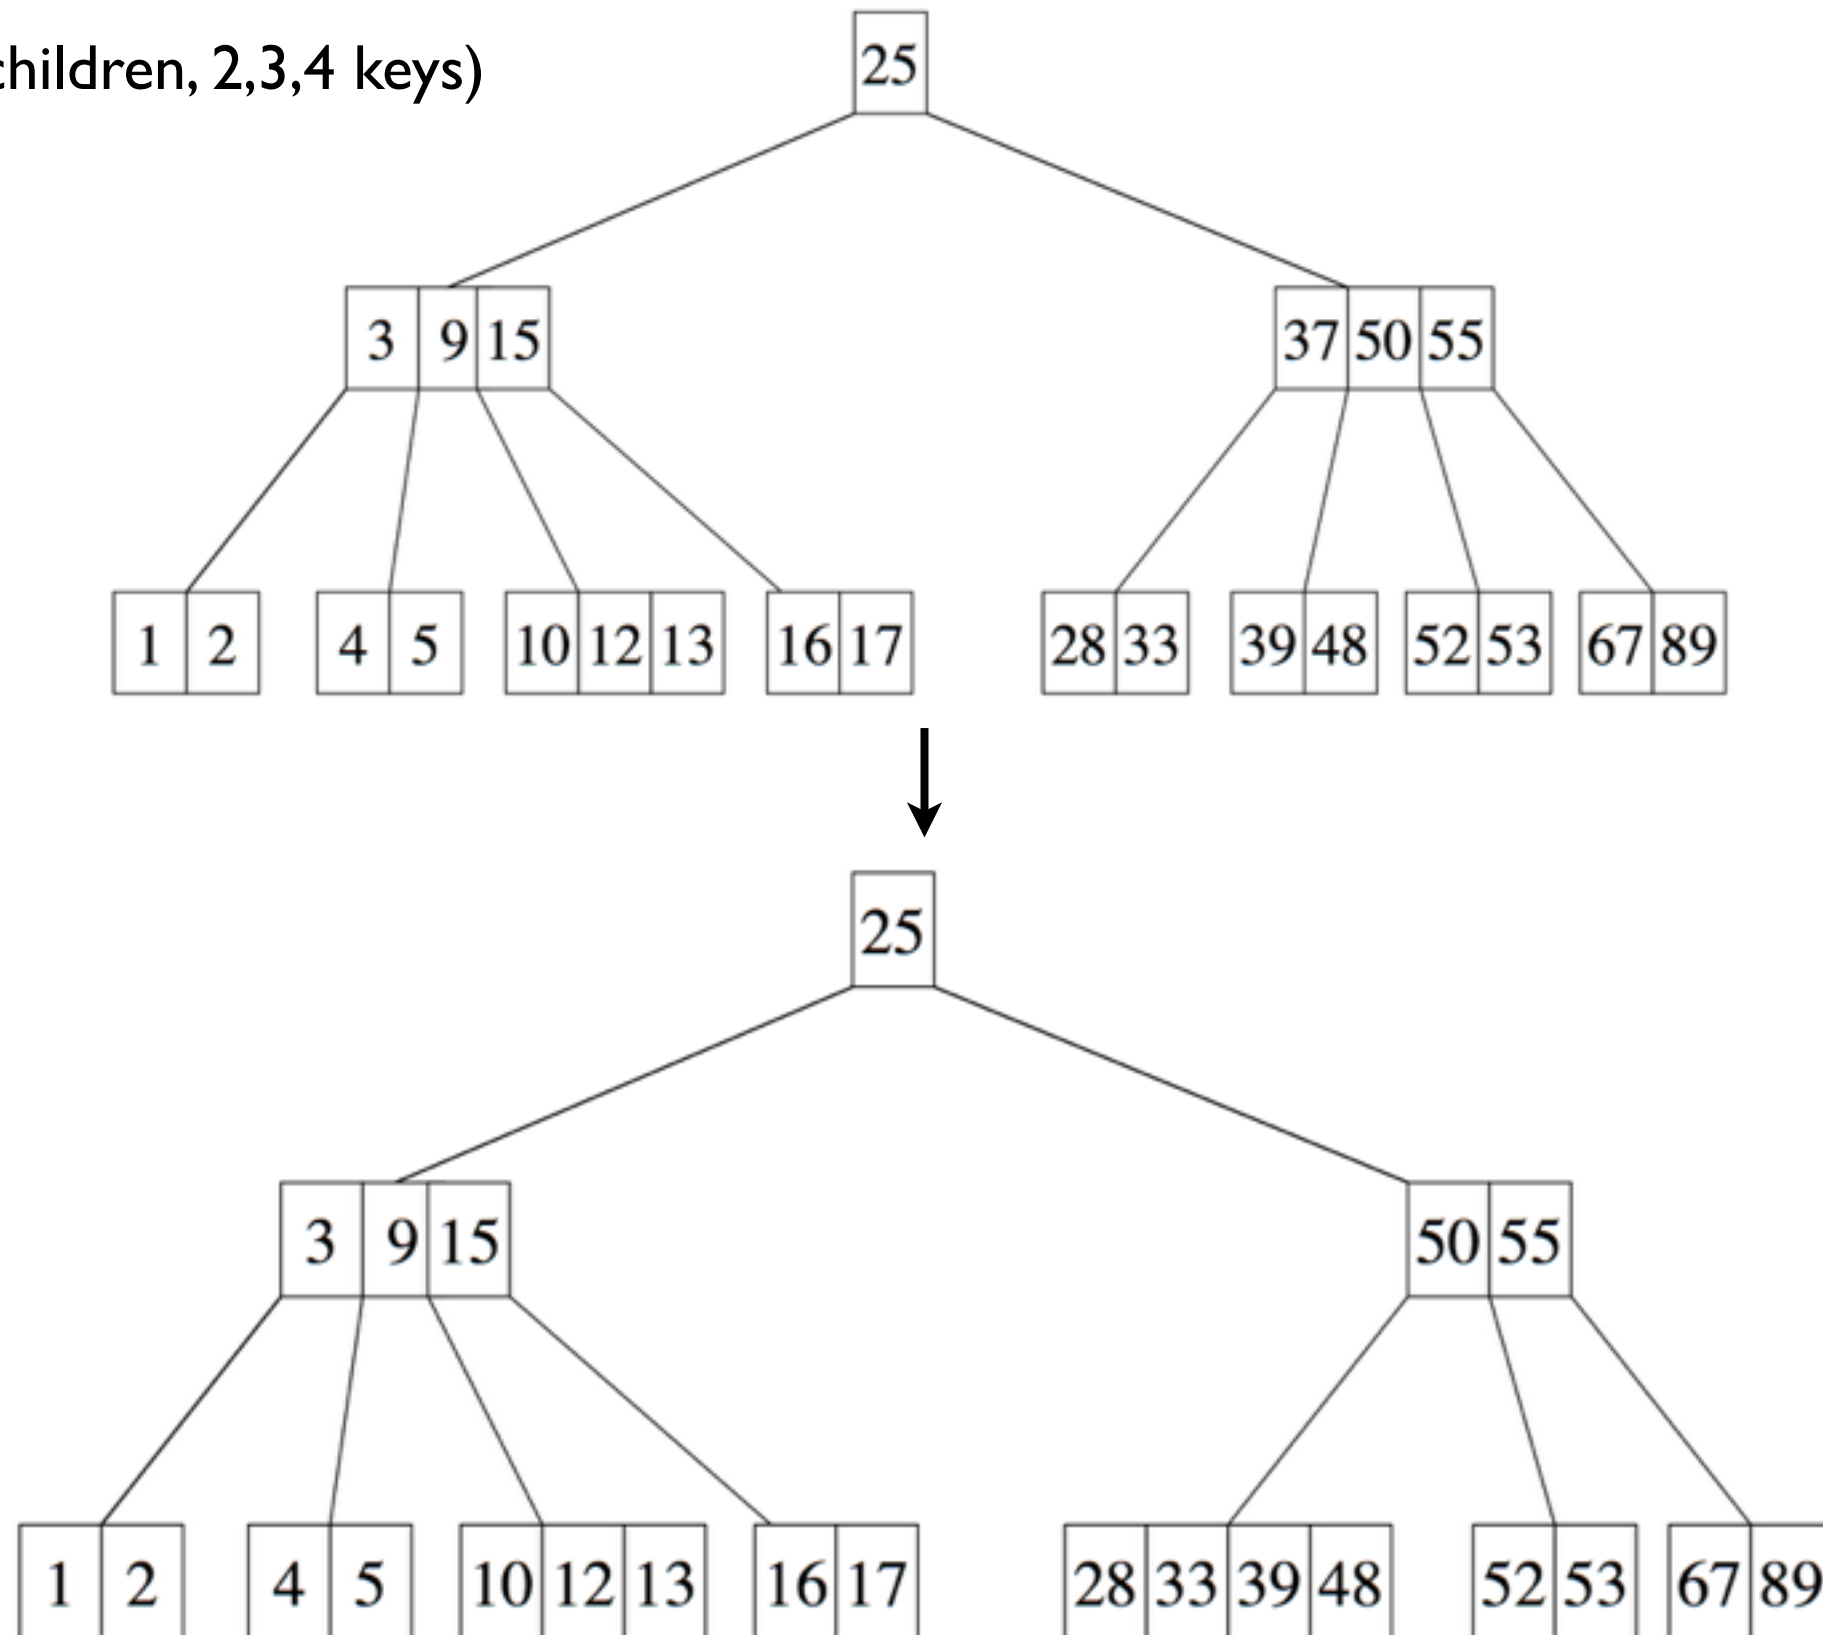
delete 37



deletion

m= 5 (3, 4, 5 children, 2,3,4 keys)

delete 37



Use of B-tree in database system

- number of disk access is $O(\log_m n)$
- each disk access requires $O(\log m)$ overhead to determine the direction to branch, but this is done in main memory without a hard disk access, thus negligible.
- m can be determined as large as possible, but it must still be small enough so that an internal node can fit into one disk block.
- m is typically between 32 and 256.
- often one or two levels of internal nodes reside in main memory.